

A1

Assignment #1 Class: Econ 613 Name: Javier Fernandez Date: March 1st, 2021

Econ 613 - Assignment 1 - Javier Fernandez

Preliminaries —

```
# Libraries
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.0 --

## v ggplot2 3.3.2      v purrr   0.3.4
## v tibble  3.0.4      v dplyr   1.0.2
## v tidyr   1.1.2      v stringr 1.4.0
## v readr   1.4.0      v forcats 0.5.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
# install.packages("matlib")
library(matlib)
```

— PART 1 —

0. Load Data —

```
dat_student <- read.csv("C:/Users/javie/OneDrive/Documents/Duke - MAE/Academic/Spring 2021/ECON 613 - Assignment 1/Student Data.csv")
dat_school <- read.csv("C:/Users/javie/OneDrive/Documents/Duke - MAE/Academic/Spring 2021/ECON 613 - Assignment 1/School Data.csv")
dat_position <- read.csv("C:/Users/javie/OneDrive/Documents/Duke - MAE/Academic/Spring 2021/ECON 613 - Assignment 1/Position Data.csv")
```

1. Missing data —

For easier calculations, I am gathering information to show one choice per row (six rows per student)

```
programs <- dat_student %>% select(1:4,11:18) %>%
  pivot_longer(cols=c("choicepgm1","choicepgm2","choicepgm3","choicepgm4",
                      "choicepgm5","choicepgm6"),names_to="Choice_num",
              names_prefix = "choicepgm",
              values_to="Program")

schools <- dat_student %>% select(1:10,17:18) %>%
  pivot_longer(cols=c("schoolcode1","schoolcode2","schoolcode3","schoolcode4",
                      "schoolcode5","schoolcode6"),names_to="Choice_num",
              names_prefix = "schoolcode",
              values_to="School")

dat_student_clean <- merge(programs,schools)
```

a. Number of students

```
nrow(dat_student) # or max(dat_student$X)
```

```
## [1] 340823
```

```
# answer: the data contains 340,823 distinct students
```

b. Number of schools

```
dat_student_clean %>% group_by(School) %>% summarise(Count=n()) %>%
  filter(School!="" | !is.na(School)) %>% nrow()
```

```
## [1] 640
```

```
# answer: the data contains 640 different schools
```

c. Number of programs

```
dat_student_clean %>% group_by(Program) %>% summarise(Count=n()) %>%
  filter(Program!="" | !is.na(Program)) %>% nrow()
```

```
## [1] 33
```

```
# answer: the data contains 33 distinct programs
```

d. Number of choices

```
# Number of distinct choices:  
dat_student_clean %>% group_by(School,Program) %>% summarise(Count=n()) %>%  
  filter(Program!="" | School!="") %>% nrow()
```

```
## [1] 3085
```

```
# answer: There are 3,085 distinct choices: combinations of schools and programs
```

e. Missing test scores

```
sum(is.na(dat_student$score))
```

```
## [1] 179887
```

```
# answer: There are 179,887 students missing test scores
```

f. Apply to the same school

```
dat_student_clean %>% group_by(X) %>% filter(Program!="") %>%  
  summarise(Number_of_schools=n_distinct(School)) %>% filter(Number_of_schools==1) %>% nrow()
```

```
## [1] 663
```

```
# answer: 663 students applied to the same school in all their cases, irrespective to the  
#          number of programs they applied to.
```

g. Apply to less than 6 choices

```
dat_student_clean %>% group_by(X) %>% filter(Program=="") %>%  
  summarise(Number_of_Programs=n()) %>% nrow()
```

```
## [1] 20988
```

```
# answer: 20,988 students applied to less than 6 choices.
```

2. Data —

```

# To do this we first have to filter the student data by only keeping the
# the student was admitted to.

program_admitted <- dat_student_clean %>% filter(Choice_num==rankplace)

admission_stats <- program_admitted %>% group_by(School,Program) %>%
  summarize(cutoff=min(score),quality=mean(score),size=n())

### erasing duplicates in dat_school
dat_school_clean <- dat_school[!duplicated(dat_school$schoolcode),]

choice_lvl_data <- left_join(admission_stats,dat_school_clean[, -1],by= c("School"="schoolcode"))

head(choice_lvl_data)

```

```

## # A tibble: 6 x 9
## # Groups:   School [1]
##   School Program  cutoff quality  size schoolname      sssdistrict ssslong ssslat
##   <int> <chr>    <int>   <dbl> <int> <chr>      <chr>      <dbl>   <dbl>
## 1  10101 Agricul~   288    310.    49 EBENEZER SENI~ Accra Metr~ -0.197    5.61
## 2  10101 Business  305    325.   100 EBENEZER SENI~ Accra Metr~ -0.197    5.61
## 3  10101 General~  316    330.   100 EBENEZER SENI~ Accra Metr~ -0.197    5.61
## 4  10101 General~  299    329.    50 EBENEZER SENI~ Accra Metr~ -0.197    5.61
## 5  10101 Home Ec~  284    301.    49 EBENEZER SENI~ Accra Metr~ -0.197    5.61
## 6  10101 Visual ~  296    312.    50 EBENEZER SENI~ Accra Metr~ -0.197    5.61

```

3. Distance —

```

# Getting the coordinates for the junior high school
program_admitted_location <- left_join(program_admitted,dat_position[, -1],by= c("jssdistrict"="jssdistrict"))

program_admitted_location <- left_join(program_admitted_location[, -1],choice_lvl_data,
  by=c("School"="School","Program"="Program"))

# Renaming variables for simplicity
program_admitted_location <- program_admitted_location %>%
  rename(jsslong=point_x,jsslat=point_y,
    sss_code=School)

# Constructing distance variable
program_admitted_location <- program_admitted_location %>%
  mutate(distance=sqrt((69.172*(ssslong-jsslong)*cos(jsslat/57.3))^2 +
    (69.172*(ssslat-jsslat))^2))

head(program_admitted_location)

```

```

##   score agey male      jssdistrict rankplace Choice_num
## 1   249   16    0      Agona Swedru         5          5
## 2   254   19    1 Abura/Asebu/Kwamankese (Abura Dunkwa) 2          2
## 3   277   17    0 Abura/Asebu/Kwamankese (Abura Dunkwa) 4          4

```

```
## 4 236 16 0 Abura/Asebu/Kwamankese (Abura Dunkwa) 3 3
## 5 237 18 1 Ajumako/Enyan/Essiam (Ajumako) 1 1
## 6 262 16 0 Twifo Hemang (Twifo Praso) 6 6
## Program sss_code jsslong jsslat cutoff quality size
## 1 General Arts 30403 -0.7552425 5.617353 208 245.2105 38
## 2 Agriculture 30403 -1.1970884 5.130001 219 241.9333 15
## 3 Home Economics 30403 -1.1970884 5.130001 215 248.3750 8
## 4 General Arts 30403 -1.1970884 5.130001 208 245.2105 38
## 5 General Arts 30403 -1.0053846 5.401725 208 245.2105 38
## 6 Agriculture 30403 -1.5597034 5.572999 219 241.9333 15
## schoolname sssdistrict
## 1 ABAKRAMPa SENIOR HIGH TECHNICAL Abura/Asebu/Kwamankese (Abura Dunkwa)
## 2 ABAKRAMPa SENIOR HIGH TECHNICAL Abura/Asebu/Kwamankese (Abura Dunkwa)
## 3 ABAKRAMPa SENIOR HIGH TECHNICAL Abura/Asebu/Kwamankese (Abura Dunkwa)
## 4 ABAKRAMPa SENIOR HIGH TECHNICAL Abura/Asebu/Kwamankese (Abura Dunkwa)
## 5 ABAKRAMPa SENIOR HIGH TECHNICAL Abura/Asebu/Kwamankese (Abura Dunkwa)
## 6 ABAKRAMPa SENIOR HIGH TECHNICAL Abura/Asebu/Kwamankese (Abura Dunkwa)
## ssslong ssslat distance
## 1 -1.197088 5.130001 45.40499
## 2 -1.197088 5.130001 0.00000
## 3 -1.197088 5.130001 0.00000
## 4 -1.197088 5.130001 0.00000
## 5 -1.197088 5.130001 22.96873
## 6 -1.197088 5.130001 39.52487
```

4. Descriptive Characteristics —

```
# Total sample
program_admitted_location %>%
  summarise(Mean_cutoff=mean(cutoff),
            Stdev_cutoff=sd(cutoff),
            Mean_quality=mean(quality),
            Stdev_quality=sd(quality),
            Mean_distance=mean(distance,na.rm = TRUE),
            Stdev_distance=sd(distance,na.rm = TRUE))

## Mean_cutoff Stdev_cutoff Mean_quality Stdev_quality Mean_distance
## 1 268.3248 52.83939 296.0099 46.02852 31.00918
## Stdev_distance
## 1 46.51059
```

By School

```
program_admitted_location %>% group_by(schoolname) %>%
  summarise(Mean_cutoff=mean(cutoff),
            Stdev_cutoff=sd(cutoff),
            Mean_quality=mean(quality),
            Stdev_quality=sd(quality),
            Mean_distance=mean(distance,na.rm = TRUE),
            Stdev_distance=sd(distance,na.rm = TRUE)) %>% head()
```

```
## # A tibble: 6 x 7
##   schoolname Mean_cutoff Stdev_cutoff Mean_quality Stdev_quality Mean_distance
##   <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 ABAKRAMP~    212.        5.48       244.        2.85       22.7
## 2 ABETIFI P~    267.        9.92       297.        6.53       45.3
## 3 ABETIFI T~    208.        9.11       247.        7.15       13.0
## 4 ABOR SENI~    210.        3.47       245.        2.36       27.7
## 5 ABUAKWA S~    324.        9.70       343.        8.29       36.9
## 6 ABURAMAN ~    204.        8.00       250.        4.78       22.3
## # ... with 1 more variable: Stdev_distance <dbl>
```

By Program

```
program_admitted_location %>% group_by(Program) %>%
  summarise(Mean_cutoff=mean(cutoff),
            Stdev_cutoff=sd(cutoff),
            Mean_quality=mean(quality),
            Stdev_quality=sd(quality),
            Mean_distance=mean(distance,na.rm = TRUE),
            Stdev_distance=sd(distance,na.rm = TRUE)) %>% head()
```

```
## # A tibble: 6 x 7
##   Program Mean_cutoff Stdev_cutoff Mean_quality Stdev_quality Mean_distance
##   <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 Accoun~    206.        13.9       245.        10.9       21.9
## 2 Agric.~    213.         5.09       256.         3.46       55.8
## 3 Agricu~    246.        39.3       275.        34.5       27.8
## 4 Auto B~    298.         0        320.         3.01       32.3
## 5 Block ~    217.        14.4       258.        16.3       30.0
## 6 Busine~    268.        52.2       298.        44.2       30.7
## # ... with 1 more variable: Stdev_distance <dbl>
```

Differentiated by quantiles (This is to be interpreted as the mean cutoff of the schools quantile x students will go to.)

```
program_admitted_location %>% mutate(quantile=case_when(
  score<quantile(score,0.1)~1,
  score>=quantile(score,0.1) & score<quantile(score,0.2)~2,
  score>=quantile(score,0.2) & score<quantile(score,0.3)~3,
  score>=quantile(score,0.3) & score<quantile(score,0.4)~4,
  score>=quantile(score,0.4) & score<quantile(score,0.5)~5,
  score>=quantile(score,0.5) & score<quantile(score,0.6)~6,
  score>=quantile(score,0.6) & score<quantile(score,0.7)~7,
  score>=quantile(score,0.7) & score<quantile(score,0.8)~8,
  score>=quantile(score,0.8) & score<quantile(score,0.9)~9,
  score>=quantile(score,0.9) ~10,
)) %>% group_by(quantile) %>% summarise(Mean_cutoff=mean(cutoff),
                                       Stdev_cutoff=sd(cutoff),
                                       Mean_quality=mean(quality),
                                       Stdev_quality=sd(quality),
```

```
Mean_distance=mean(distance,na.rm = TRUE),
Stdev_distance=sd(distance,na.rm = TRUE))
```

```
## # A tibble: 10 x 7
##   quantile Mean_cutoff Stdev_cutoff Mean_quality Stdev_quality Mean_distance
##   <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1      1      210.        9.19       246.        10.7       25.1
## 2      2      220.       14.2       254.        11.5       26.0
## 3      3      229.       17.9       261.        13.1       26.9
## 4      4      239.       21.3       269.        15.3       27.7
## 5      5      251.       24.3       279.        18.1       29.6
## 6      6      265.       26.4       292.        20.4       30.8
## 7      7      278.       27.7       304.        21.8       31.3
## 8      8      295.       30.2       319.        24.0       32.9
## 9      9      324.       29.7       345.        24.7       35.5
## 10     10      366.       29.5       385.        27.0       43.7
## # ... with 1 more variable: Stdev_distance <dbl>
```

— PART 2 —

5. Data creation —

```
set.seed(123)

# X1
X1 <- runif(10000,min=1,max=3)

# X2
X2 <- rgamma(10000,shape = 3,scale = 2)

# X3
X3 <- rbinom(10000,size=1,prob = 0.3)

# Error term
error <- rnorm(10000,mean=2,sd=1)

# Create Y and Ydum

# Y
par <- c(0.5,1.2,-0.9,0.1)
Y <- par[1] + par[2]*X1 + par[3]*X2 + par[4]*X3 + error

# Ydum
Ydum <- as.numeric(Y>mean(Y))
```

6. OLS —

```
# Correlation Y and X1. How different is it from 1.2?
# answer: the result is 0.21.
#Being Y a linear function of X1 we would have expected the correlation to be larger.
```

```
cor(Y,X1)
```

```
## [1] 0.216015
```

```
# Calculate the coefficients
```

```
regressors <- as.matrix(t(rbind(rep(1,10000),X1,X2,X3)),ncol=4)
```

```
betas <-inv(t(regressors)%*%regressors)%*%(t(regressors)%*%Y)
# betas are the coefficients for the OLS estimation
```

```
resids <- Y-regressors%*%betas
sigma_2 <- as.numeric(t(resids)%*%resids/(10000-4))
var_cov_matrix <- sigma_2*inv(t(regressors)%*%regressors)
std_errors <- sqrt(diag(var_cov_matrix))
# std_errors are the standard errors of the coefficients
coef_stderrors <- cbind(betas,std_errors)
colnames(coef_stderrors) <- c("Coefs","Std. Errors")
```

```
print(coef_stderrors) # Answer
```

```
##           Coefs Std. Errors
## [1,]  2.49051092 0.040620582
## [2,]  1.19777741 0.017358659
## [3,] -0.89640329 0.002875798
## [4,]  0.08781299 0.021694686
```

7. Discrete choice —

```
# The linear probability model can be estimated by OLS
# Function:
linear_prob_model <- function(Y,regressors){
  betas <-inv(t(regressors)%*%regressors)%*%(t(regressors)%*%Y)
  resids <- Y-regressors%*%betas
  sigma_2 <- as.numeric(t(resids)%*%resids/(nrow(regressors)-ncol(regressors)))
  var_cov_matrix <- sigma_2*inv(t(regressors)%*%regressors)
  std_errors <- sqrt(diag(var_cov_matrix))
  coef_stderrors <- cbind(betas,std_errors)
  colnames(coef_stderrors) <- c("Coefs","Std. Errors")
  return(coef_stderrors) # Answer
}

# Estimation
regressors <- as.matrix(t(rbind(rep(1,10000),X1,X2,X3)),ncol=4)
```



```
# Results of the linear probability model (Coefs and regressors)
results_lpm <- linear_prob_model(Ydum,regressors)
```

```
#### 7.b Probit ----
```

```
# Probit function
probit_likelihood = function(coefs,x1,x2,x3,y)
{
  xbeta      = coefs[1] + coefs[2]*x1 + coefs[3]*x2 + coefs[4]*x3
  pr         = pnorm(xbeta)
  pr[pr>0.999999] = 0.999999
  pr[pr<0.000001] = 0.000001
  like       = y*log(pr) + (1-y)*log(1-pr)
  return(-sum(like))
}
```

```
### Estimation
```

```
# Result Probit----
start = runif(4)
res_probit = optim(start,fn=probit_likelihood,method="BFGS",control=list(trace=6,REPORT=10,maxit=10000)
```

7.a. Linear probability model —

```
## initial value 24689.277344
## iter 10 value 2214.628584
## final value 2213.313307
## converged
```

```
fisher_info_probit = solve(res_probit$hessian)
prop_sigma_probit = sqrt(diag(fisher_info_probit))
```

```
#### 7.c Logit ----
```

```
# Logit function
logit_likelihood = function(y,x1,x2,x3,coefs)
{
  xbeta      = coefs[1] + coefs[2]*x1 + coefs[3]*x2 + coefs[4]*x3
  pr         = exp(xbeta)/(1+exp(xbeta))
  pr[pr>0.999999] = 0.999999
  pr[pr<0.000001] = 0.000001
  like       = y*log(pr) + (1-y)*log(1-pr)
  return(-sum(like))
}
```

```
### Estimation
```

```
# Result Logit ----
start = runif(4)
res_logit = optim(start,fn=logit_likelihood,method="BFGS",control=list(trace=6,REPORT=10,maxit=10000),x
```

```
## initial value 20564.984166
## iter 10 value 2224.622518
## final value 2223.017344
## converged
```

```
fisher_info_logit = solve(res_logit$hessian)
prop_sigma_logit = sqrt(diag(fisher_info_logit))
```

```
# Final Results ----
results = cbind(par,results_lpm[,1],results_lpm[,2],
                res_probit$par,prop_sigma_probit,res_logit$par,prop_sigma_logit)
colnames(results) = c("True parameter","LPM: est","LPM :se","Probit: est","Probit: :se",
                      "Logit: est","Logit: :se")
results
```

```
##      True parameter      LPM: est      LPM :se Probit: est Probit: :se
## [1,]           0.5 0.885860391 0.0136557488 3.04275799 0.10007791
## [2,]           1.2 0.146150735 0.0058356006 1.17235964 0.04292123
## [3,]          -0.9 -0.102941654 0.0009667803 -0.90546589 0.01858996
## [4,]           0.1 -0.008099353 0.0072932778 -0.01124976 0.04647615
##      Logit: est Logit: :se
## [1,] 5.42655537 0.18557806
## [2,] 2.10059552 0.07936241
## [3,] -1.61851052 0.03670961
## [4,] -0.01963215 0.08323293
```

```
# Answer
# 1. The LPM, which is the only comparable model in terms of coefficients, produces estimates
# really different from the true parameters. At least they are all un the correct direction.
# Using a t-test with 95% confidence, the intercept, X1 and X2 are significant.
# This is not the case for X3.
Significance_lpm <- abs(results_lpm[,1]/results_lpm[,2])>1.96
Significance_lpm
```

```
## [1] TRUE TRUE TRUE FALSE
```

```
# 2. The Probit model coefficients are not directly comparable with the true parameters.
# We can observe that the sign of the estimates is correct for all but X3.
# Using a t-test with 95% confidence, the intercept, X1 and X2 are significant.
# This is not the case for X3.
Significance_probit <-abs(res_probit$par/prop_sigma_probit)>1.96
Significance_probit
```

```
## [1] TRUE TRUE TRUE FALSE
```

```
# 3. The Logit model coefficients are not directly comparable with the true parameters.
# We can observe that the sign of the estimates is correct for all but X3.
# Using a t-test with 95% confidence, the intercept, X1 and X2 are significant.
# This is not the case for X3.
Significance_logit <- abs(res_logit$par/prop_sigma_logit)>1.96
Significance_logit
```

```
## [1] TRUE TRUE TRUE FALSE
```

8. Marginal effects —

```
#### 8.1. Probit average marginal effects
Xbeta_probit <- regressors %*% as.matrix(res_probit$par)
mgl_effects_probit <- pnorm(Xbeta_probit)%*% t(as.matrix(res_probit$par))

mean_mgleff_probit <- colMeans(mgl_effects_probit)
sd_mgleff_probit <- apply(mgl_effects_probit,2,sd)

#### 8.2. Logit average marginal effects
Xbeta_logit <- regressors %*% as.matrix(res_logit$par)
mgl_effects_logit <- (plogis(Xbeta_logit)*(1-plogis(Xbeta_logit)))*%*%t(as.matrix(res_logit$par))

mean_mgleff_logit <- colMeans(mgl_effects_logit)
sd_mgleff_logit <- apply(mgl_effects_logit,2,sd)

## Answers

avg_mgleffects <- cbind(mean_mgleff_probit,sd_mgleff_probit,mean_mgleff_logit,sd_mgleff_logit)
avg_mgleffects <- avg_mgleffects[-1,]
colnames(avg_mgleffects) <- c("Probit: Avg Mgl Eff","Probit: SD of Mgl Eff",
                             "Logit: Avg Mgl Eff","Logit: SD of Mgl Eff")

rownames(avg_mgleffects) <- c("X1","X2","X3")
avg_mgleffects # Answer
```

```
##      Probit: Avg Mgl Eff Probit: SD of Mgl Eff Logit: Avg Mgl Eff
## X1      0.658724865      0.492616340      0.14403075
## X2     -0.508762736      0.380469678     -0.11097581
## X3     -0.006321007      0.004727059     -0.00134611
##      Logit: SD of Mgl Eff
## X1      0.175445743
## X2      0.135181085
## X3      0.001639715
```

```
# The average marginal effects are, in general, larger in the probit model than in the logit model.
# The same is true for the standard errors. In both models, X1 has the largest marginal effect.
```